MULTI-AGENT DELIBERATION

RUST · CARGO

CONVERGED

# Think Tank

Four agents debate. An Arbiter decides when they have converged — measured, not guessed.

ENGINEERING REPORT · RENDERED FROM REPOSITORY DATA

HOW IT WORKS

# A strictly sequential debate loop

- Researcher claims, Skeptic challenges, Visionary improvises, Synthesizer merges — then the Arbiter scores.
- The Arbiter weighs semantic stabilisation at 0.80 and the opposition ratio at 0.20.
- Rounds repeat until the alignment score clears the target, or the round budget runs out.

| 4 | 0.80 | 6 |
| --- | --- | --- |
| AGENTS PER ROUND | SEMANTIC WEIGHT | ROUND CEILING |

PROVEN ON ITS OWN OUTPUT

# This deck is the measurement

- Rendered by deckgen: a headless egui harness painting ArtCraft's tokens onto A4.
- One binary, no browser, no network — any repo's data can produce this deck.




**0** BROWSERS IN THE LOOP

ENGINEERING

# Boring on purpose

ONE RUST CRATE, NO REDUCERS

every node clones state and appends, so history is never lost.

All providers speak the OpenAI protocol; adding one is a row in a table.

Tests drive a local stub server — the suite makes zero network calls.

| 5 | 3 | 0 |
| --- | --- | --- |
| NODES PER ROUND | STRUCTURED-OUTPUT RETRIES | NETWORK CALLS IN TESTS |

THE HONEST NUMBERS

# What the arbiter actually saw

## 1.00

CONVERGED SCORE - AGAINST A 0.65 TARGET, IN 2 OF 6 ROUNDS

THE FORCED RUN

The same loop scored 0.1536 against an impossible 1.01 target.

A benchmark that only shows wins is marketing — this one shows both.

| 0.15 | 392 |
| --- | --- |
| FORCED SCORE | TOKENS SPENT |

WHERE THIS GOES

# Decks as a build artifact

## Any repo that emits JSON can emit a deck

Reports, benchmarks and architecture reviews render from the same file. A coding agent regenerates this PDF from repository data — no browser, no network, no layout code.

**1**
DESIGN SYSTEM

**3**
OFL TYPEFACES

**6**
PAGES FROM ONE FILE